



CMSC 206: Database Management Systems

MongoDB

Thomas LAURENT

21/11/2020

University of the Philippines Open University

The Student Evaluation of Teachers (SET) survey is still available. Please do fill it in, I am always eager to know what I can do better (but also what I did well :)).

Introduction

This week we will introduce a new DBMS, MongoDB¹. This will mark a break from relational databases, and our first touch with NoSQL.

We will first introduce the document-oriented data model, on which MongoDB is based and then survey the MongoDB equivalents of the SQL commands we saw previously.

¹<https://www.mongodb.com>

Document-oriented databases

Document databases - Introduction

Document-oriented databases represent data as *documents* each given given a *unique identifier*. Related documents are grouped in *collections*. A big difference with the relational model is that *no schema is specified for the data*. The documents are seen as self-describing data.

You can think of documents a bit like objects in object oriented programming, with different attributes that have values. As there is no schema though, two objects in the same collection might not have the same attributes. Documents can be stored in different ways, as XML, YAML, or JSON objects for example. MongoDB uses JSON.

Another difference is that an attribute's value can be a document itself, or an array. Attributes are no longer simple values as in the classic relational model we studied before.

Document databases - Example

Relational

ID	first_name	last_name	cell	city	year_of_birth	location_x	location_y
1	'Mary'	'Jones'	'516-555-2048'	'Long Island'	1986	'-73.9876'	'40.7574'

ID	user_id	profession
10	1	'Developer'
11	1	'Engineer'

ID	user_id	name	version
20	1	'MyApp'	1.0.4
21	1	'DocFinder'	2.5.7

ID	user_id	make	year
30	1	'Bentley'	1973
31	1	'Rolls Royce'	1965

MongoDB

```
{
  first_name: "Mary",
  last_name: "Jones",
  cell: "516-555-2048",
  city: "Long Island",
  year_of_birth: 1986,
  location: {
    type: "Point",
    coordinates: [-73.9876, 40.7574]
  },
  profession: ["Developer", "Engineer"],
  apps: [
    { name: "MyApp",
      version: 1.0.4 },
    { name: "DocFinder",
      version: 2.5.7 }
  ],
  cars: [
    { make: "Bentley",
      year: 1973 },
    { make: "Rolls Royce",
      year: 1965 }
  ]
}
```

Figure 1: Example of relational model to document model in MongoDB from <https://www.mongodb.com/document-databases>. This shows a document that could be part of People collection in a database for example.

Data modelling document-oriented database - Embedding

As we see in the previous example, a way to represent relationships between entities is to embed them into one another, e.g. embedding car documents into the person document in figure 1. Embedding usually leads to better read speed and lets us easily retrieve related data: when we retrieve the “parent” document we also get all the related, embedded data.

Embedding is particularly adapted to one-to-one or one-to-many relationships. When used for many-to-many relationships it leads to data redundancy: the “child” document must be repeated in each “parent”.

Embedding can also lead to uncontrolled growth of the “parent” document - if we have many embedded documents - which could hinder performance, or even break the DBMS document size limit.

Data modelling document-oriented database - Referencing

Even in a document database we can still use the idea of “foreign keys” to represent relationships.

A big difference with foreign keys in relational databases is that here we have no constraints to ensure the integrity of the data. Ensuring the integrity is left to the application logic instead of being done by the DBMS.

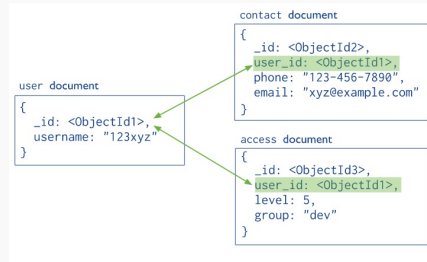


Figure 2: Example of reference between documents in 3 different collections

Data modelling document-oriented database - Embedding VS Referencing

Choosing when to use embedding or linking can be a complex problem, and depends on many factors. We should consider the performance trade off we are willing to make (e.g. can we accept some data redundancy for better read performance with embedding?), and the data we are modelling.

MongoDB has a lot of information available on the subject, as well as examples and patterns *here*

MongoDB

Introduction

MongoDB is a very popular² document-oriented, JSON based database. It is designed to be used as a distributed database with replication and high availability and uses an ad-hoc query language that integrates javascript.

The documentation for MongoDB is available *here*

To use mongoDB, you can either *install it on your system* or use the hosted *here*, creating a free cluster, and *connect to your database*.

In the next slides we will see the basics of MongoDB usage. As usual, the documentation should be your first stop for more complete and complex information.

²see <https://db-engines.com/en/ranking>

MongoDB does not enforce a schema on the data, so there is no need for a DDL. That said, schema validation is possible, see *here*

Selecting the database to use

```
use dbName
```

If the database does not exist, it will be created when data is inserted in it. Similarly, when inserting a document in a collection, the collection will be created if it does not already exist. That said, we can also explicitly create a collection and specify options for it, such as a maximum number of documents in the collection, or schema validation rules.

Creating a collection

```
db.createCollection(name, options)
```

```
db.createCollection("myCollection",  
    {"capped": true, "max": 100})
```

DML - insert

Inserting data into a database is very straightforward, as we just have to specify the collection we want to insert data in and the document(s) we want to insert.

Inserting data

`db.collection.insert(document or array of documents)`

```
db.people.insert({  
  FName: "Thomas",  
  LName: "Laurent",  
  Nationality: "French"  
})
```

```
db.carBrands.insert([  
  {name: "Renault"},  
  {name: "Volkswagen"}  
])
```

To update a collection, we have to specify a *query* (WHERE clause in SQL) and the *update* we want to apply (SET clause in SQL).

We can also specify options on how the update should behave, e.g the *multi* option that decides if the update should update only one record matching the query or all of them, with the default being only updating one.

Updating data

```
db.collection.update(query, update, options)
db.collection.updateOne(query, update, options)
db.collection.updateMany(query, update, options)
```

Updating data

```
db.collection.update(query, update, options)
```

The *query* is expressed with the same operators we will describe later in the DQL. The update can be expressed in different ways:

- By specifying a new document that should replace the updated document(s)
- By specifying an *update document* that contains *update operators*³ such as **\$set** to set an attribute in the updated documents or **\$inc** to increase an attribute's value.

³see

<https://docs.mongodb.com/manual/reference/operator/update/#id1>

DML - Example with new document

```
db.people.update(  
  { name: "Andy"},    // Query parameter  
  {                   // Update document  
    name: "Andy",  
    rating: 1,  
    score: 1  
  }  
)
```

This will replace one document where the name is "Andy" with the update document

DML - Example with new document

```
db.books.update(  
  { _id: 1 },  
  {  
    $inc: { stock: 5 },  
    $set: {  
      item: "ABC123",  
      "info.publisher": "2222",  
      tags: [ "software" ],  
      "ratings.1": { by: "xyz", rating: 3 }  
    }  
  }  
)
```

This will set (“\$set”) the “item”, “info.publisher”, “tags”, and “ratings.1” attributes to the given values; and increase (or increment “\$inc”) the value of stock by 5 in the document with id number 1.

Deleting data

```
db.collection.deleteOne(query) db.collection.deleteMany(query)
```

```
db.orders.deleteOne(  
  { "expiryts" :  
    { $lt: ISODate("2015-11-01T12:40:15Z") }  
  }  
);
```

This will delete one document in collection “orders” with an expiryts before (less than, “\$lt”) the 1st of November 2015 at 12:50:15

Querying data

```
db.collection.find(query) db.collection.findOne(query)
```

Mongo does not use SQL, it uses its own ad-hoc query language language based on operators. The main operators are query and projection operators, detailed *here*, and the lookup operator, that lets us "join" collections, detailed *here*

A good, simple, interactive introduction to queries in MongoDB, with comparisons with SQL is given *here*

By default MongoDB has no access control, so no need for DCL.

Still, if need be, it can be activated and the managing of users, roles and permissions is detailed *here*